

Automatic Differentiation

Import autograd and create a variable

```
In [1]: from mxnet import autograd, nd  
  
x = nd.arange(4).reshape((4, 1))  
print(x)
```

```
[[0.]  
 [1.]  
 [2.]  
 [3.]]  
<NDArray 4x1 @cpu(0)>
```

Attach gradient to \mathbf{x}

- It allocates memory to store its gradient, which has the same shape as \mathbf{x} .
- It also tell the system that we need to compute its gradient.

```
In [2]: x.attach_grad()  
x.grad
```

```
Out[2]: [[0.]  
         [0.]  
         [0.]  
         [0.]]  
<NDArray 4x1 @cpu(0)>
```

Forward

Now compute

$$y = 2\mathbf{x}^T \mathbf{x}$$

by placing code inside a `with autograd.record():` block. MXNet will build the according computation graph.

```
In [3]: with autograd.record():  
        y = 2 * nd.dot(x.T, x)  
        y
```

```
Out[3]: [[28.]]  
        <NDArray 1x1 @cpu(0)>
```

Backward

In [4]: `y.backward()`

Get the gradient

Given $y = 2\mathbf{x}^\top \mathbf{x}$, we know

$$\frac{\partial y}{\partial \mathbf{x}} = 4\mathbf{x}$$

Now verify the result:

```
In [5]: print((x.grad - 4 * x).norm().asscalar() == 0)
        print(x.grad)
```

True

```
[[ 0.]
 [ 4.]
 [ 8.]
 [12.]]
<NDArray 4x1 @cpu(0)>
```

Backward on non-scalar

`y.backward()` equals to `y.sum().backward()`

```
In [6]: with autograd.record():  
        y = 2 * x * x  
        print(y.shape)  
        y.backward()  
        print(x.grad)
```

```
(4L, 1L)
```

```
[[ 0.]
```

```
 [ 4.]
```

```
 [ 8.]
```

```
 [12.]]
```

```
<NDArray 4x1 @cpu(0)>
```

Training Mode and Prediction Mode

The `record` scope will alter the mode by assuming that gradient is only required for training.

It's necessary since some layers, e.g. batch normalization, behavior differently in the training and prediction modes.

```
In [7]: print(autograd.is_training())  
with autograd.record():  
    print(autograd.is_training())
```

False

True

Computing the Gradient of Python Control Flow

Autograd also works with Python functions and control flows.

```
In [8]: def f(a):  
        b = a * 2  
        while b.norm().asscalar() < 1000:  
            b = b * 2  
        if b.sum().asscalar() > 0:  
            c = b  
        else:  
            c = 100 * b  
        return c
```


Function behaviors depends on inputs

```
In [9]: a = nd.random.normal(shape=1)
        a.attach_grad()
        with autograd.record():
            d = f(a)
        d.backward()
```

Verify the results

f is piecewise linear in its input a . There exists g such as $f(a) = ga$ and $\frac{\partial f}{\partial a} = g$. Verify the result:

```
In [10]: print(a.grad == (d / a))
```

```
[1.]  
<NDArray 1 @cpu(0)>
```

Head gradients and the chain rule

We can break the chain rule manually. Assume $\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x}$. `y.backward()` will only compute $\frac{\partial y}{\partial x}$. To get $\frac{\partial z}{\partial x}$, we can first compute $\frac{\partial z}{\partial y}$, and then pass it as head gradient to `y.backward`.

```
In [11]: with autograd.record():
          y = x * 2
          y.attach_grad()
          with autograd.record():
              z = y * x
              z.backward() # y.grad = \partial z / \partial y
              y.backward(y.grad)
              x.grad == 2*x # x.grad = \partial z / \partial x
```

```
Out[11]: [[1.]
          [1.]
          [1.]
          [1.]]
          <NDArray 4x1 @cpu(0)>
```