


Neural Style Transfer


```
In [ ]: %matplotlib inline
import d2l
from mxnet import autograd, gluon, image, init, nd
from mxnet.gluon import model_zoo, nn
import time
```

Read the Content and Style Images

```
In [8]: d2l.set_figsize()  
content_img = image.imread('rainier.jpg')  
d2l.plt.imshow(content_img.asnumpy());
```



```
In [9]: style_img = image.imread('autumn_oak.jpg')
d2l.plt.imshow(style_img.asnumpy());
```



Features

Preprocessing and Postprocessing

```
In [3]: rgb_mean = nd.array([0.485, 0.456, 0.406])
        rgb_std = nd.array([0.229, 0.224, 0.225])

        def preprocess(img, image_shape):
            img = image.imresize(img, *image_shape)
            img = (img.astype('float32') / 255 - rgb_mean) / rgb_std
            return img.transpose((2, 0, 1)).expand_dims(axis=0)

        def postprocess(img):
            img = img[0].as_in_context(rgb_std.context)
            return (img.transpose((1, 2, 0)) * rgb_std + rgb_mean).clip(0, 1)
```

Select Content and Style Layers

```
In [5]: pretrained_net = model_zoo.vision.vgg19(pretrained=True)
style_layers, content_layers = [0, 5, 10, 19, 28], [25]
net = nn.Sequential()
for i in range(max(content_layers + style_layers) + 1):
    net.add(pretrained_net.features[i])
```

Extract Features

```
In [7]: def extract_features(X, content_layers, style_layers):  
        contents = []  
        styles = []  
        for i in range(len(net)):  
            X = net[i](X)  
            if i in style_layers:  
                styles.append(X)  
            if i in content_layers:  
                contents.append(X)  
        return contents, styles
```

Get Contents and Style Features

```
In [8]: def get_contents(image_shape, ctx):  
        content_X = preprocess(content_img, image_shape).copyto(ctx)  
        contents_Y, _ = extract_features(content_X, content_layers, style_layers)  
        return content_X, contents_Y  
  
        def get_styles(image_shape, ctx):  
            style_X = preprocess(style_img, image_shape).copyto(ctx)  
            _, styles_Y = extract_features(style_X, content_layers, style_layers)  
            return style_X, styles_Y
```

Define the Loss Function

Content Loss

```
In [9]: def content_loss(Y_hat, Y):  
        return (Y_hat - Y).square().mean()
```


Style Loss

```
In [10]: def gram(X):  
    num_channels, n = X.shape[1], X.size // X.shape[1]  
    X = X.reshape((num_channels, n))  
    return nd.dot(X, X.T) / (num_channels * n)  
  
    def style_loss(Y_hat, gram_Y):  
        return (gram(Y_hat) - gram_Y).square().mean()
```

Total Variance Loss

```
In [12]: def tv_loss(Y_hat):  
    return 0.5 * ((Y_hat[:, :, 1:, :] - Y_hat[:, :, :-1, :]).abs().mean() +  
                  (Y_hat[:, :, :, 1:] - Y_hat[:, :, :, :-1]).abs().mean())
```

Loss Function

```
In [13]: style_channels = [net[l].weight.shape[0] for l in style_layers]
content_weight, style_weight, tv_weight = 1, 1e3, 10

def compute_loss(X, contents_Y_hat, styles_Y_hat, contents_Y, styles_Y_gram):
    contents_l = [content_loss(Y_hat, Y) * content_weight for Y_hat, Y in zip(
        contents_Y_hat, contents_Y)]
    styles_l = [style_loss(Y_hat, Y) * style_weight for Y_hat, Y in zip(
        styles_Y_hat, styles_Y_gram)]
    tv_l = tv_loss(X) * tv_weight
    l = nd.add_n(*styles_l) + nd.add_n(*contents_l) + tv_l
    return contents_l, styles_l, tv_l, l
```

Create and Initialize the Composite Image

```
In [14]: class GeneratedImage(nn.Block):  
    def __init__(self, img_shape, **kwargs):  
        super(GeneratedImage, self).__init__(**kwargs)  
        self.weight = self.params.get('weight', shape=img_shape)  
  
    def forward(self):  
        return self.weight.data()
```

Initialize

```
In [15]: def get_inits(X, ctx, lr, styles_Y):  
    gen_img = GeneratedImage(X.shape)  
    gen_img.initialize(init.Constant(X), ctx=ctx, force_reinit=True)  
    trainer = gluon.Trainer(gen_img.collect_params(), 'adam',  
                           {'learning_rate': lr})  
    styles_Y_gram = [gram(Y) for Y in styles_Y]  
    return gen_img(), styles_Y_gram, trainer
```

Train

```
In [16]: def train(X, contents_Y, styles_Y, ctx, lr, max_epochs, lr_decay_epoch):
        X, styles_Y_gram, trainer = get_inits(X, ctx, lr, styles_Y)
        for i in range(max_epochs):
            start = time.time()
            with autograd.record():
                contents_Y_hat, styles_Y_hat = extract_features(
                    X, content_layers, style_layers)
                contents_l, styles_l, tv_l, l = compute_loss(
                    X, contents_Y_hat, styles_Y_hat, contents_Y, styles_Y_gram)
            l.backward()
            trainer.step(1)
            nd.waitall()
            if i % 50 == 0 and i != 0:
                print('epoch %3d, content loss %.2f, style loss %.2f, '
                      'TV loss %.2f, %.2f sec'
                      % (i, nd.add_n(*contents_l).asscalar(),
                         nd.add_n(*styles_l).asscalar(), tv_l.asscalar(),
                         time.time() - start))
            if i % lr_decay_epoch == 0 and i != 0:
                trainer.set_learning_rate(trainer.learning_rate * 0.1)
                print('change lr to %.1e' % trainer.learning_rate)
        return X
```

Train a 150 x 225 composite image

```
In [17]: ctx, image_shape = d2l.try_gpu(), (225, 150)
net.collect_params().reset_ctx(ctx)
content_X, contents_Y = get_contents(image_shape, ctx)
style_X, styles_Y = get_styles(image_shape, ctx)
output = train(content_X, contents_Y, styles_Y, ctx, 0.01, 500, 200)
d2l.plt.imsave('neural-style-1.png', postprocess(output).asnumpy())

epoch 50, content loss 10.12, style loss 29.39, TV loss 3.46, 0.03 sec
epoch 100, content loss 7.50, style loss 15.44, TV loss 3.90, 0.03 sec
epoch 150, content loss 6.31, style loss 10.38, TV loss 4.15, 0.03 sec
epoch 200, content loss 5.66, style loss 8.10, TV loss 4.29, 0.03 sec
change lr to 1.0e-03
epoch 250, content loss 5.60, style loss 7.93, TV loss 4.30, 0.03 sec
epoch 300, content loss 5.54, style loss 7.78, TV loss 4.31, 0.03 sec
epoch 350, content loss 5.48, style loss 7.63, TV loss 4.31, 0.03 sec
epoch 400, content loss 5.42, style loss 7.48, TV loss 4.32, 0.03 sec
change lr to 1.0e-04
epoch 450, content loss 5.41, style loss 7.47, TV loss 4.32, 0.03 sec
```

Visualize



Train a 300 x 450 composite image

```
In [19]: image_shape = (450, 300)
content_X, content_Y = get_contents(image_shape, ctx)
style_X, style_Y = get_styles(image_shape, ctx)
X = preprocess(postprocess(output) * 255, image_shape)
output = train(X, content_Y, style_Y, ctx, 0.01, 300, 100)
d2l.plt.imsave('neural-style-2.png', postprocess(output).asnumpy())

epoch 50, content loss 13.89, style loss 13.71, TV loss 2.38, 0.10 sec
epoch 100, content loss 9.54, style loss 8.79, TV loss 2.65, 0.10 sec
change lr to 1.0e-03
epoch 150, content loss 9.23, style loss 8.46, TV loss 2.68, 0.10 sec
epoch 200, content loss 8.95, style loss 8.17, TV loss 2.70, 0.10 sec
change lr to 1.0e-04
epoch 250, content loss 8.92, style loss 8.14, TV loss 2.70, 0.10 sec
```

Visualize

